

1. SQL Basics

October 25, 2013

Directions

- Work in groups of 2 students.
- Write-up a short note (max 2 pages long) to sum up your experiments and focus on singular difficulties of the subject.
- Please, add at the end of your report, the following sentence:

“We, ⟨student1⟩ and ⟨student2⟩, attest this is our own work and accept the consequences if it is not.”

- Deliverables must be packaged into a *tar.gz* or *zip* file containing a single directory having your short report (PDF file) and the SQL scripts to play back the entire story. Name of the directory must satisfy the pattern:

`hw1_<name of student1>_<name of student2>`

- The archive must be uploaded on the Moodle web site up to the due date.

Requirements

PostgreSQL (≥ 9) database system is required to achieve this assignment. It is an Open Source software available for download at <http://www.postgresql.org/download/>. The abundant documentation and active community could help you installing and setting the system if necessary. The *PgAdmin III* client included into the regular install package is the preferred way for the connection to the database.

Due date

2013-10-11 11:59pm

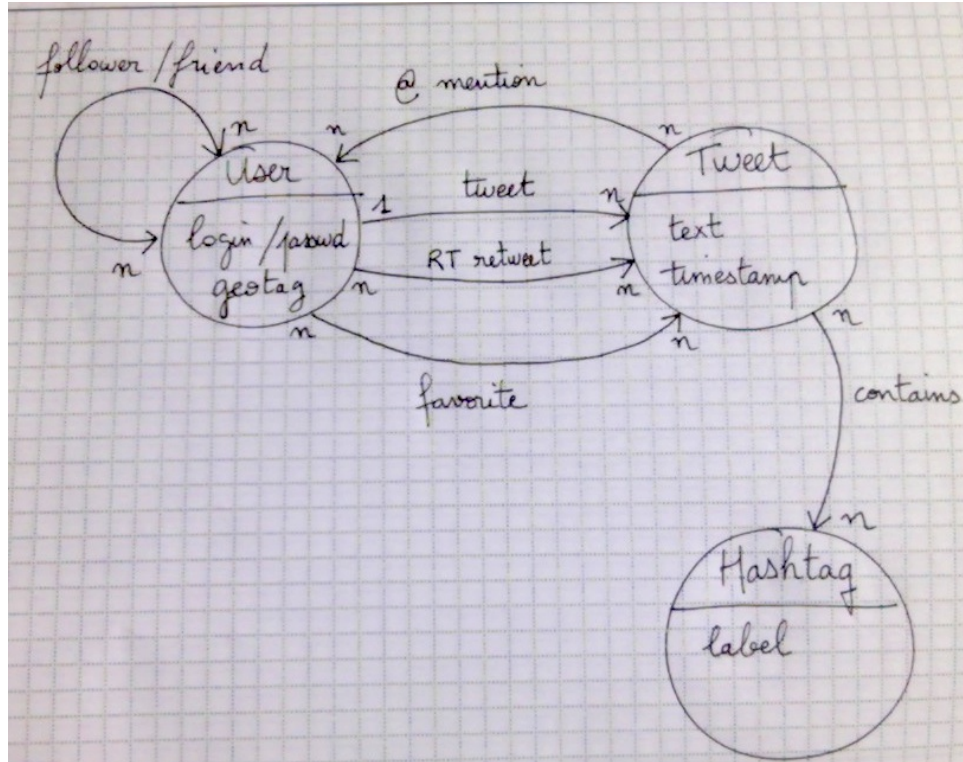


Figure 1: *TinyTwitter* — A (Not So) Simplified Model of Twitter

Introduction

The playground is Twitter, the famous micro-blogging platform. We are trying to develop a restricted copy of Twitter, coined *TinyTwitter*.

The drawing on Figure gives a model for *TinyTwitter*. There are three main entities (user, tweet, hashtag) and several binary relationships among them. A registered user (with username/password) post tweets, i.e. short messages of less than 80 characters each. Tweets may actually be retweets (RT <tweetid>) of previous messages from registered users. RT is idempotent, i.e. RT RT <tweetid> = RT <tweetid>. Users might follow any other registered users, so-called friends. Friend's tweets are highlighted. Tweets admit hashtags and mentions. Hashtags (#Nantes) are basically keywords for tweet categorization purpose. Mentions such like @me are references to registered users and allow for instance to answer a previous tweet or to focus on your audience.

The above model translates into the set of following relations:

User	(uid, username, password, geotag)
Tweet	(tid, text, timestamp, user)
Hashtag	(hid, label)
Tagintweet	(tag, tweet)
Mention	(tweet, user)
Retweet	(tweet, retweet)
Follower	(user, friend)
Favorite	(user, tweet)

1 DDL part of SQL and more

1. Write SQL `CREATE TABLE` statements to implement the eight relations of the *TinyTwitter* project. Propose accurate types for every field and incorporate integrity constraints such like:
 - *primary keys*, *foreign keys*;
 - field is *not null*;
 - field is *unique*;
 - domain-based and tuple-based constraints (*check*);
2. Write SQL `INSERT` statements to populate the database. The resulting dataset is required to return a non-empty answer set to any query below.

2 DML part of SQL

Write SQL statements to translate the following queries against the above *TinyTwitter* database.

1. (a) List all tweets of user **me** in descending timestamp order.
(b) List the 10 more recent tweets from friends of **me**.
(c) Give location of users that tweet about **#Nantes**.
2. (a) Give the average number of followers by user.
(b) Give the number of tweets a day for the entire last month.
3. (a) List in descending timestamp order both tweets that mention user **me**, and retweets from **me**.
(b) Find users that never mention their friends.
(c) Find users that have retweeted all their friends.

Resources & References

1. Chapter *Relational Model* of Lecture Notes (and SQL counterparts)
2. Basic SQL tutorials and quick reference guide
<http://www.w3schools.com/sql/>
3. Homepage of the PostgreSQL 9.3 documentation
<http://www.postgresql.org/docs/9.3/static/index.html>
4. Data integrity PostgreSQL document page
<http://www.postgresql.org/docs/9.3/static/ddl-constraints.html>